
Machine Learning Final Project: Sepsis Prediction

Team Members:

Akshay Jambhulkar(23PGAI004)

Avinash Gupta(23PGAI0085)

Syed Nizamuddin(23PGAI0060)

Vishnuvardhan C(23PGAI0100)

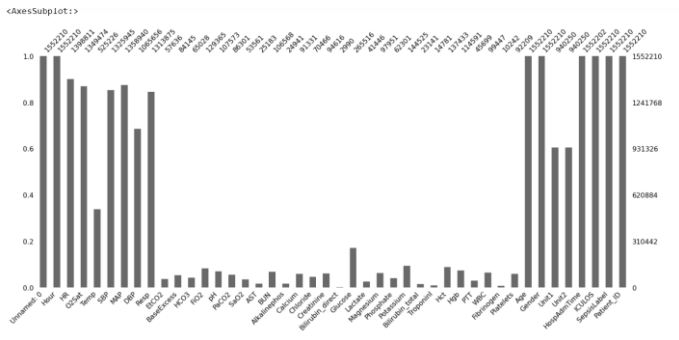
Introduction

Objective:

The main goal of this project is to use Machine learning algorithms (Supervised & Unsupervised) to predict sepsis diagnosis 6 hours prior to sepsis manifestation using training SetA, training SetB. These datasets are hospital 1 & hospital 2.

Data Preprocessing:

On data preprocessing we find out that many features have missing values, & more than 90% of values are missing for multiple columns in the dataset.



O2Sat	0.130611
Temp	0.661627
SBP	0.145770
MAP	0.124513
DBP	0.313459
Resp	0.153546
EtcCO2	0.962868
BaseExcess	0.945790
HC03	0.958106
FiO2	0.916658
pH	0.930697
PaCO2	0.944401
SaO2	0.965494
AST	0.983776
BUN	0.931344
Alkalinephos	0.983932
Calcium	0.941161
Chloride	0.954603
Creatinine	0.939044
Bilirubin_direct	0.998074
Glucose	0.828943
Lactate	0.973299
Magnesium	0.936896
Phosphate	0.959863
Potassium	0.906891
Bilirubin_total	0.985092
TroponinI	0.990477
Hct	0.911460
Hgb	0.926176
PTT	0.970559
WBC	0.935932
Fibrinogen	0.983402

Missing values percentage



Introduction

Objective:

The main goal of this project is to use Machine learning algorithms (Supervised & Unsupervised) to predict sepsis diagnosis 6 hrs prior to sepsis manifestation using training SetA, training SetB. These datasets are hospital 1 & hospital 2.

Data Set:

- We are using physiological data which is hosted on PhysioNet Computing in Cardiology Challenge 2019.
- Retrieving the dataset which is having pipe separated features for respective timestamp.
- Data contains 40 feature consisting of:
 - 8 Vitals Signs – Heart Rate, Temperature, MAP, ...
 - 28 Laboratory Values – FiO2, Lactate, Bilirubin, ...
 - 6 Demographics – Age, Gender, Hospital Unit, ...

Example of one patient dataset post imputation:

	HR	O2Sat	SBP	MAP	DBP	Resp	Hour	Age	Gender	HospAdmTime	ICULOS	Identifier	SepsisLabel
0	96.537964	98.533285	114.147204	69.647189	55.645112	20.816360	0.0	27.92	1.0	-0.03	1.0	9.0	0.0
1	117.000000	99.000000	116.000000	97.000000	81.000000	20.000000	1.0	27.92	1.0	-0.03	2.0	9.0	0.0
2	96.793996	98.526412	114.230461	69.827433	55.800597	20.861704	2.0	27.92	1.0	-0.03	3.0	9.0	0.0
3	96.922011	98.522975	114.272089	69.917555	55.878340	20.884377	3.0	27.92	1.0	-0.03	4.0	9.0	0.0
4	97.050027	98.519539	114.313718	70.007677	55.956082	20.907049	4.0	27.92	1.0	-0.03	5.0	9.0	0.0

Data Preprocessing:

- On data preprocessing we find out that many features have missing values, & more than 90% of values are missing for multiple columns in the dataset.

Data imputation & preparing training , test & validation files

Data imputation & Data split into train, test & validation:

- Linear imputation (forward & backward fill) with nearest neighbor.
- Imputed each patient separately before merging the .psv files as a single file.
- Patients (.psv files) in Training Set A was divided in three different sets: Training – 10%, Validation Set -15% and Test- 15%

Baseline Function for imputing missing values

```
#function to fill the missing values, we will use backward & ffill method to fill the missing values for each patient,
# since the data is provided patient wise psv files, we will also impute the data patient wise using backward & forward fill

def impute_missing_vals(df, features):
    df_clean=df.copy()
    for feature in features:
        if df_clean[feature].isnull().sum()!=len(df_clean):
            df_clean[feature]=df_clean[feature].fillna(0)
        elif df_clean[feature].isnull().sum()!=len(df_clean)-1:
            df_clean[feature]=df_clean[feature].ffill().bfill()
        else:
            df_clean[feature]=df_clean[feature].interpolate(method='nearest',limit_direction='both')
            df_clean[feature]=df_clean[feature].ffill().bfill()
    return df_clean
```

```
In [5]: #After imputation of missing values the columns which get the least unique values need to be dropped,hence we will check unique
# value count & drop those columns

columns_to_drop=[]
for column in train_df.columns:
    vc=len(train_df[column].value_counts().unique())
    if vc<5:
        columns_to_drop.append(column)
    else:
        pass
print(columns_to_drop)
['EtCO2', 'Gender', 'Unit1', 'Unit2', 'Sepsislabel']

In [6]: #here we cannot drop ['Gender', 'Sepsislabel']
items_to_remove=['Gender', 'Sepsislabel']
for column in columns_to_drop:
    if column in items_to_remove:
        columns_to_drop.remove(column)
    else:
        pass
print(columns_to_drop)
['EtCO2', 'Unit1', 'Unit2']
```

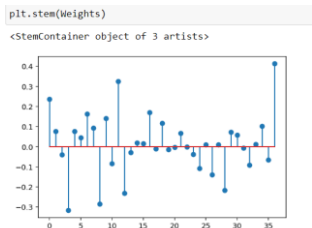
Dropped above features

- Above same procedure was applied for Training Set B.
- After data imputation we looked at the features with continuous data which have unique values less than 5 & drop these features. These features were: EtCO2, Unit1, Unit2

Model Application-Logistic Regression

Logistic Regression for Classification-1:

- Logistic Regression Classifier was applied & the weight coefficients of individual features were interpreted. The 4 most important features were observed to be [ICULOS,PaCO2,SBP,HCO3]. These 4 features contribute most for the prediction of Sepsis disease.



```
from sklearn.linear_model import LogisticRegression
# We need to select a range of values for which we will need to run these
>linear_model.LogisticRegression()
penal ('l1', 'l2')
(scp, logspace(-1, 1, 10))

param_grid = {'penalty': ('l1', 'l2')}

from sklearn.model_selection import GridSearchCV
logistic_class_best = GridSearchCV(LogisticRegression, param_grid, scoring='accuracy')
logistic_class_best.fit(X_train, y_train)

out[0]: GridSearchCV(cv=...,
                    estimator=LogisticRegression(C=10, max_iter=1000, solver='saga'),
                    param_grid={'C': array([ 0.1, ..., 10.00000000, 1.
                    2.15848902, 0.04198853, 10., ..., 21.44434049, 0.00010000, 1.
                    0.00000000, 100., ..., 1.]),
                    penalty': ('l1', 'l2')},
                    scoring='accuracy')

In [21]: best_model_logistic_class_best.fit(X_train, y_train)

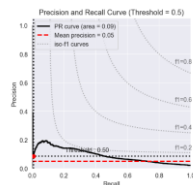
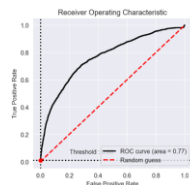
In [22]: best_model_logistic_class_best.parameters

out[22]: {'C': 0.1, 'penalty': 'l1'}
```

Hyperparametric Tuning & Performance Metrics:

Hyperparametric tuning was done & the best parameters were found. i.e {'C':0.1, 'penalty': 'l1'}. The model was run using the obtained hyper-parameters. The performance of the model was evaluated based on precision, recall, f1 score and roc curve area.

Classification Report				
	precision	recall	f1-score	support
nonsepsis	0.98	1.00	0.99	115683
Sepsis	0.09	0.01	0.02	2635
accuracy	0.98			118318
macro avg	0.53	0.50	0.50	118318
weighted avg	0.96	0.98	0.97	118318

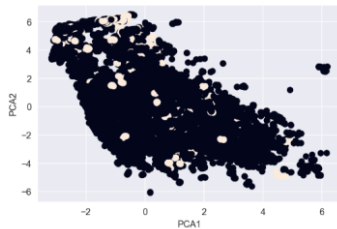


Observations: The ROC Curve (Area) for the Logistic Regression model is 0.77, Precision is 9%.

Model Application 2- PCA Along with Logistic Regression

Logistic Regression for Classification along with PCA:

PCA was applied to the pre-processed data, and the scatter plot for 1st and 2nd principal component was plotted. Post PCA, hyper parameter tuning was done to find the optimal number of principal components which gives the max information about the original features. Number of principal components to be selected were found out to be 15.



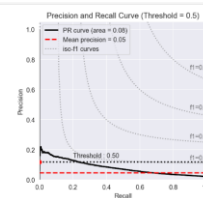
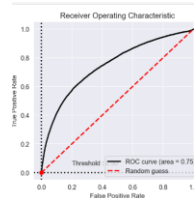
PCA 1 vs PCA 2

Performance metrics:

Hyperparametric tuning was done & the best parameters were found. i.e {'n_components=15}. The model was run using the obtained hyper-parameters. The performance of the model was evaluated based on precision, recall, f1 score and roc curve area.

```
bc_pca_lr.print_report()
```

	precision	recall	f1-score	support
nonSepsis	0.98	1.00	0.99	657478
Sepsis	0.12	0.00	0.00	14511
accuracy			0.98	671989
macro avg	0.55	0.50	0.50	671989
weighted avg	0.96	0.98	0.97	671989

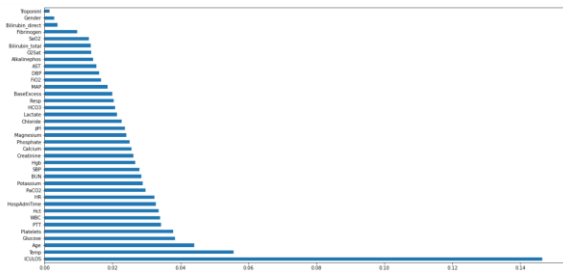


Observations: The ROC Curve (Area) for the PCA Logistic Regression model is 0.75, Precision is 12%.

Model Application 3- Random Forest

Random Forest for Classification:

Ensemble learning model (Random Forest) was applied to the dataset. The feature importance were found out and plotted as below. The feature importance was decided based on the features capability to detect sepsis.

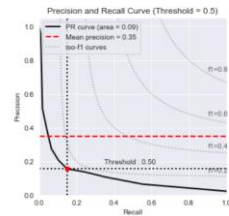
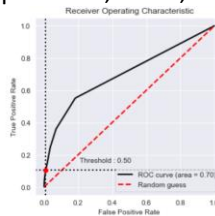


Performance metrics :

The model was run and the performance of the model was evaluated based on precision, recall, f1 score and roc curve area.

```
: bc.print_report()
```

Classification Report				
	precision	recall	f1-score	support
nonSepsis	0.98	0.99	0.99	115683
Sepsis	0.21	0.10	0.14	2635
accuracy			0.97	118318
macro avg	0.59	0.55	0.56	118318
weighted avg	0.96	0.97	0.97	118318



Observations: The ROC Curve (Area) for the Random Forest model is 0.70, Precision is 21%.

Model Application 4- xgboost

Xgboost for Classification:

Another ensembling model XGBoost was applied to the dataset.

Performance metrics :

The model was run and the performance of the model was evaluated based on precision, recall, f1 score and roc curve area

	precision	recall	f1-score	support
0	0.98	1.00	0.99	115683
1	0.22	0.04	0.07	2635
accuracy			0.98	118318
macro avg	0.60	0.52	0.53	118318
weighted avg	0.96	0.98	0.97	118318

f1score: 0.07318611987381704
Precision: 0.21682242990654205
Recall: 0.04402277039848197
accuracy: 0.9751686133977924
auc: 0.5202004017358107

Observations: The ROC Curve (Area) for the xgboost model is 0.52, Precision is 22%.

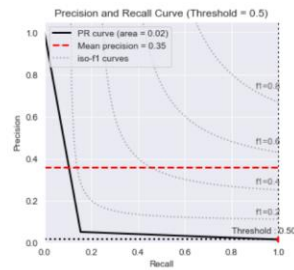
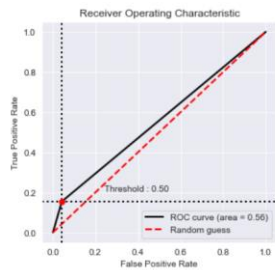
Results:

Results: After analyzing the 4 models , we are proceeding with Random Forest Classifier for the final deployment model & will test the Random Forest Classifier on the TrainingSetB (i.e. validation set)

```
In [32]: print(classification_report(y_validation_setB,y_pred_setB))
print('f1score:', f1_score(y_validation_setB,y_pred_setB))
print('Precision:', precision_score(y_validation_setB,y_pred_setB))
print('Recall:', recall_score(y_validation_setB,y_pred_setB))
print('accuracy:', accuracy_score(y_validation_setB,y_pred_setB))
print('auc:',auc)
```

	precision	recall	f1-score	support
0	0.99	0.96	0.97	751239
1	0.05	0.15	0.08	10780
accuracy			0.95	762019
macro avg	0.52	0.56	0.52	762019
weighted avg	0.97	0.95	0.96	762019

f1score: 0.07579311758409439
Precision: 0.050227897008663105
Recall: 0.15435992578849722
accuracy: 0.9467454223582351
auc: 0.5562378925278274



Observations: The ROC Curve (Area) for the Random Forest model is 0.52, Precision is 22%.

Conclusion:

Random Forest has highest diagnostic accuracy among ensemble method.

